

Database

By
Jitendra Singh Tomar || Jeetu

Topics to discuss

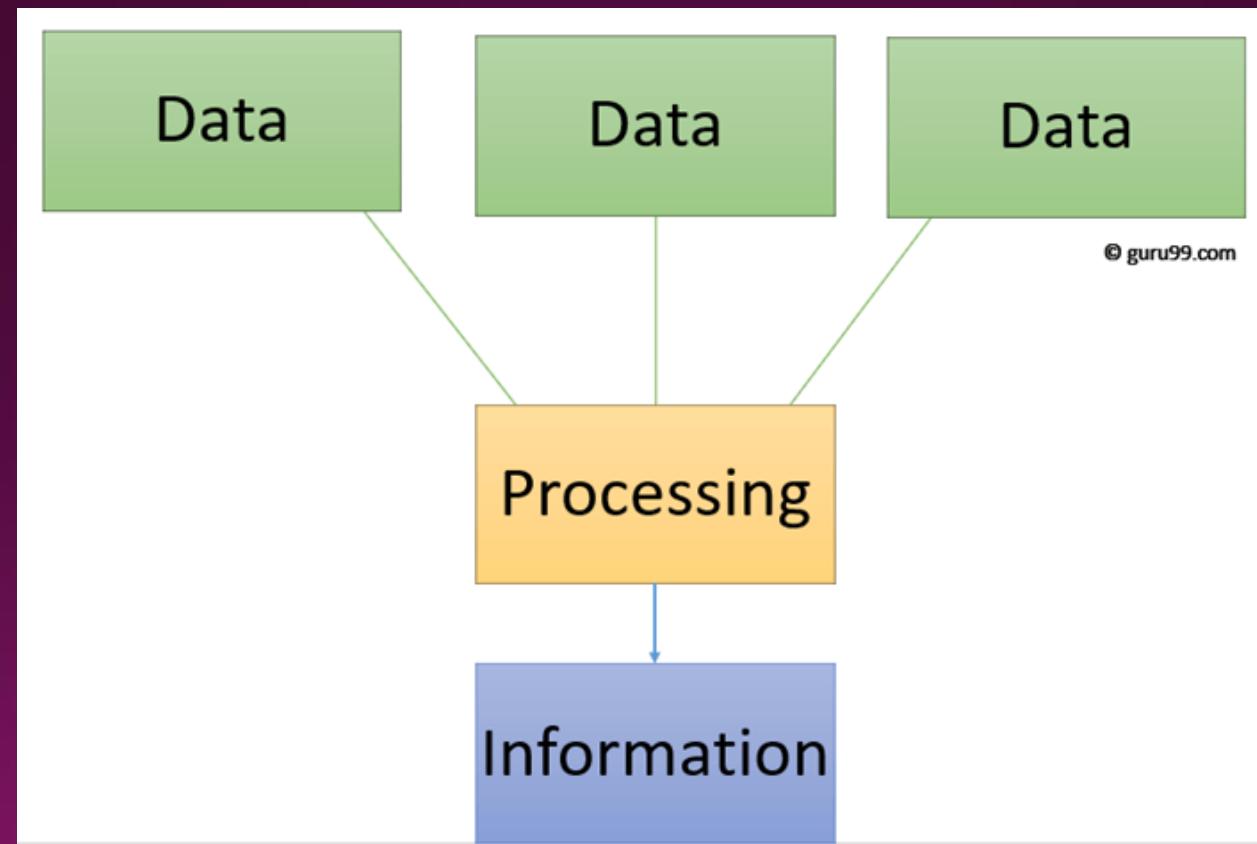
1. Understanding Databases
2. Overview of Relational Databases
3. Getting Started with SQL
4. Types of databases & their usage
5. Database Design and Modeling
6. Normalization
7. Practical Database Design
8. Advanced SQL and Database Administration Basics
9. Database Administration Basics
10. Data Integrity and Transactions

What is DATA?

- Data is raw fact or figures or entity.
- When activities in the organization takes place, the effect of these activities need to be recorded which is known as Data.

What is INFORMATION?

- Processed data is called information.
- The purpose of data processing is to generate the information required for carrying out the business activities.



Introduction to Databases

- A database is a structured collection of data that is organized and stored in a way that allows for efficient retrieval, management, and update of information.
- Databases are fundamental to modern computing and play a crucial role in storing and managing vast amounts of data for various applications and industries.

Database

- Database may be defined in simple terms as a collection of data
- A database is a collection of related data.
- The database can be of any size and of varying complexity.
- A database may be generated and maintained manually or it may be computerized.

Database Management System

- Database Management Systems (DBMS) are software systems used to store, retrieve, and run queries on data.
- A DBMS serves as an interface between an end-user and a database, allowing users to create, read, update, and delete data in the database.
- A DBMS is a collection of program that enables user to create and maintain a database.

Characteristics of DBMS

- Data Definition and Schema Management
- Data Manipulation
- Transaction Management
- Query Optimization
- Security and Access Control
- Security and Access Control
- Data Recovery and Backup
- Scalability
- Data Independence
- Query Language Support
- Backup and Recovery

DBMS Utilities

- A data loading utility:
 - Which allows easy loading of data from the external format without writing programs.
- A backup utility:
 - Which allows to make copies of the database periodically to help in cases of crashes and disasters.
- Recovery utility:
 - Which allows to reconstruct the correct state of database from the backup and history of transactions.

DBMS Utilities

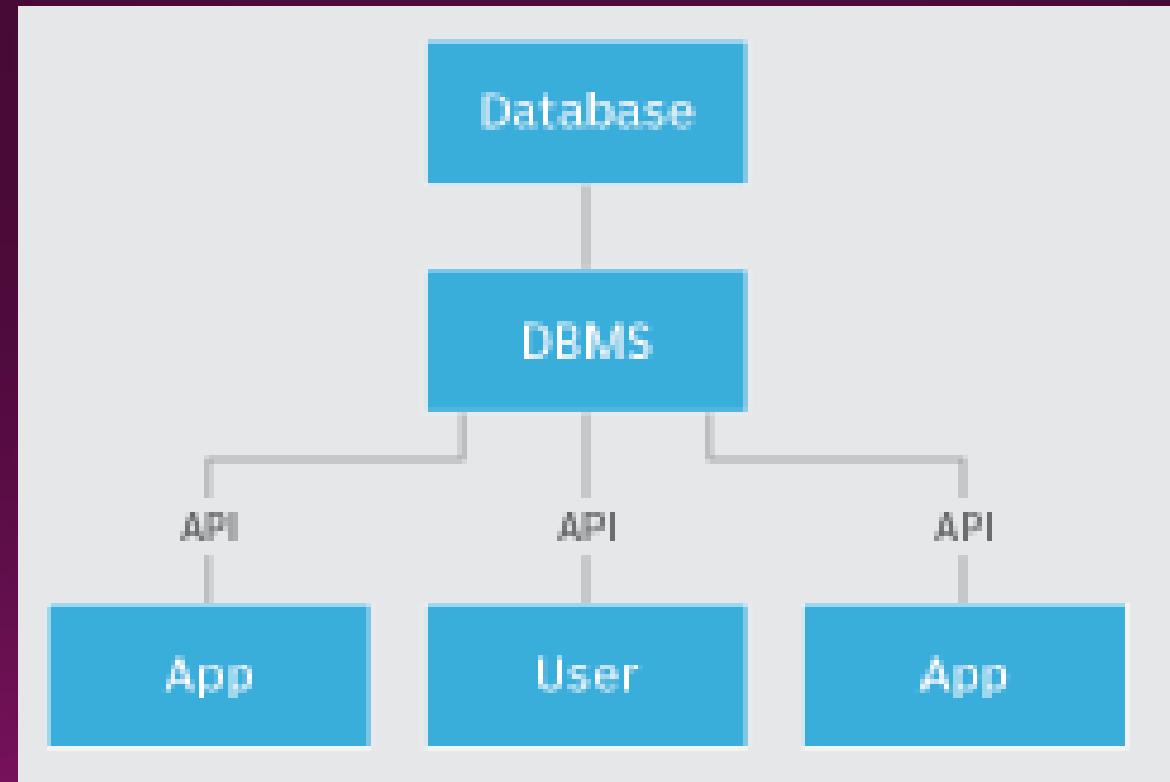
- Monitoring tools:
 - Which monitors the performance so that internal schema can be changed and database access can be optimized.
- File organization:
 - Which allows restructuring the data from one type to another?

File system VS Database

File System	DBMS
File system is a collection of data. Any management with the file system, user has to write the procedures.	DBMS is a collection of data and user is not required to write the procedures for managing the database.
File system gives the details of the data representation and Storage of data.	DBMS provides an abstract view of data that hides the details.
In File system storing and retrieving of data cannot be done efficiently.	DBMS is efficient to use since there are wide varieties of sophisticated techniques to store and retrieve the data.
Concurrent access to the data in the file system has many problems.	DBMS takes care of Concurrent access using some form of locking.
File system doesn't provide crash recovery mechanism.	DBMS has crash recovery mechanism, DBMS protects user from the effects of system failures.
Protecting a file under file system is very difficult.	DBMS has a good protection mechanism.

Advantages of DBMS

- Data independency
- Efficient data access
- Data integrity and security
- Data Administration
- Concurrent access and Crash recovery
- Reduced application development time



Types of DBMS

Relational database

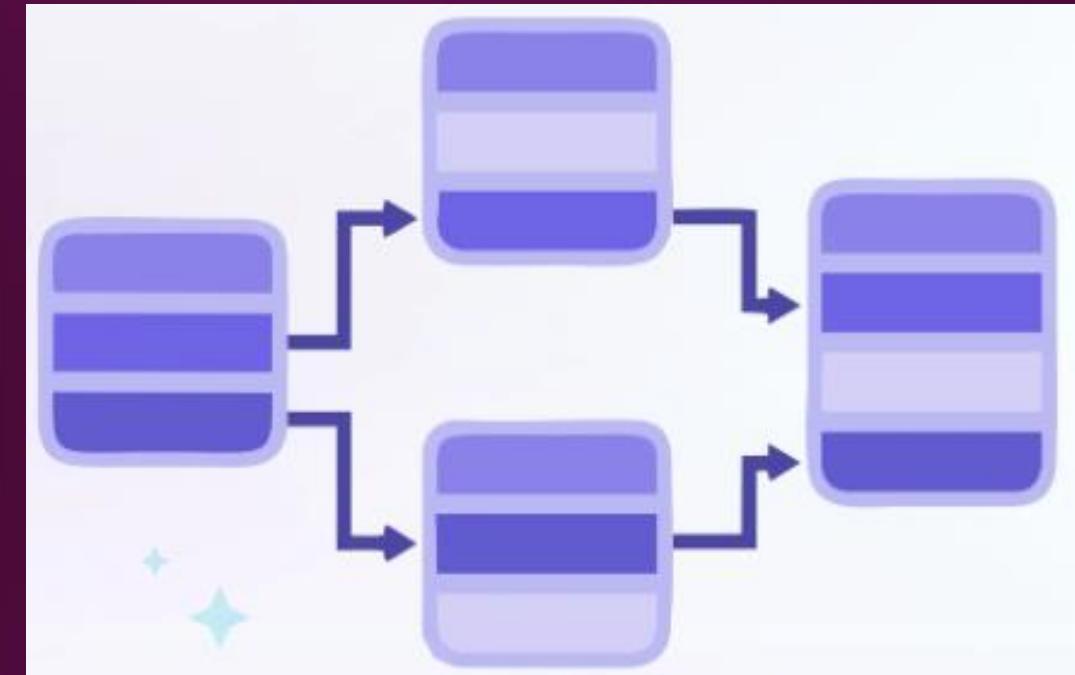
Network database

Object oriented database

Hierarchical database

Relation Database

- It is based on SQL.
- A relational database management system.
- This is one of the most popular data models which is used in industries.
- Every table in a database has a key field which uniquely identifies each record.
- RDBMS is a system where data is organized in two-dimensional tables using rows and columns.



Object Oriented Database

- It is a combination of relational database concepts and object-oriented principles.
- OOPs principles are data encapsulation, inheritance, and polymorphism.
- It requires less code and is easy to maintain.
- For example – Object DB software.

**Object-Oriented
Programming**

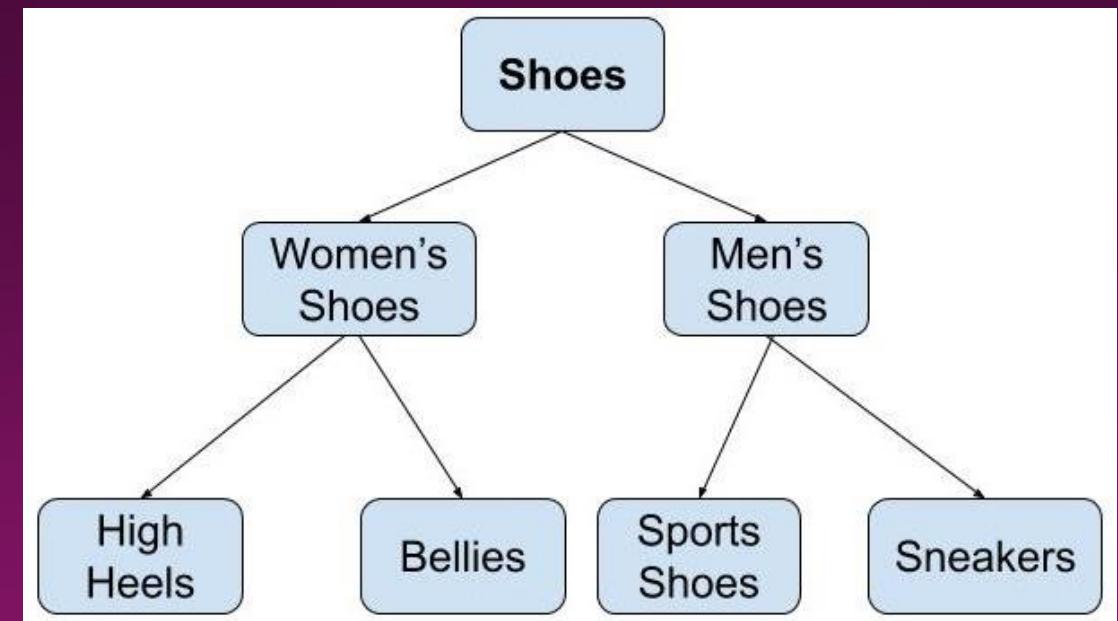
Polymorphism

Inheritance

Encapsulation

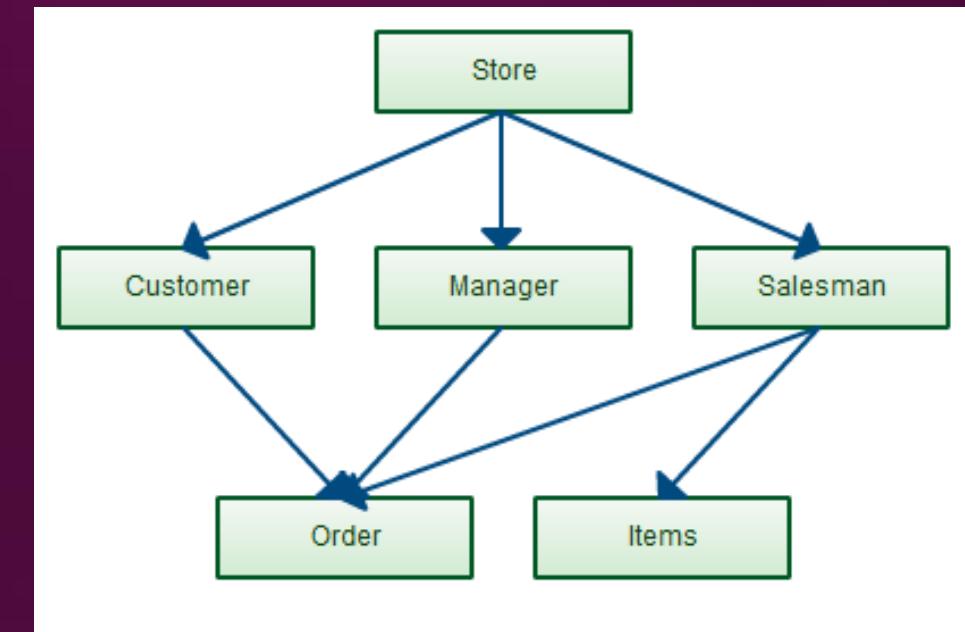
Hierarchical Database

- It is used in industry on mainframe platforms.
- The hierarchy starts from the root node, connecting all the child nodes to the parent node.
- For example
 - IMS (IBM)
 - Windows registry (Microsoft).

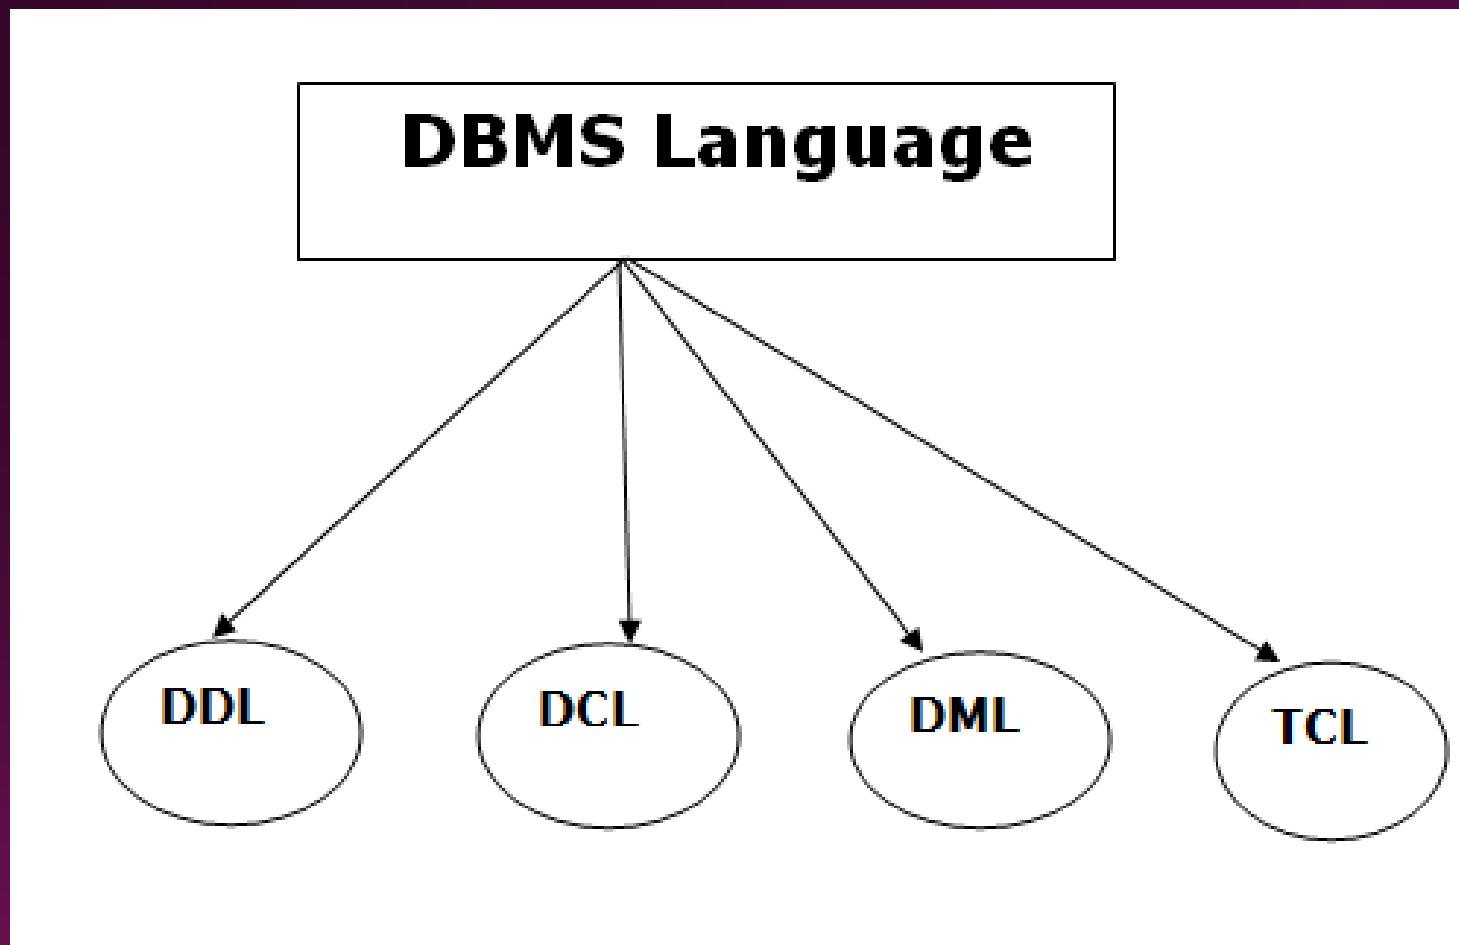


Network database

- A Network database management system.
- This maintain **one to one relationship** (1: 1) or **many to many relationship** (N: N).
- It is based on a network data model, which allows each record to be related to multiple primary records and multiple secondary records.

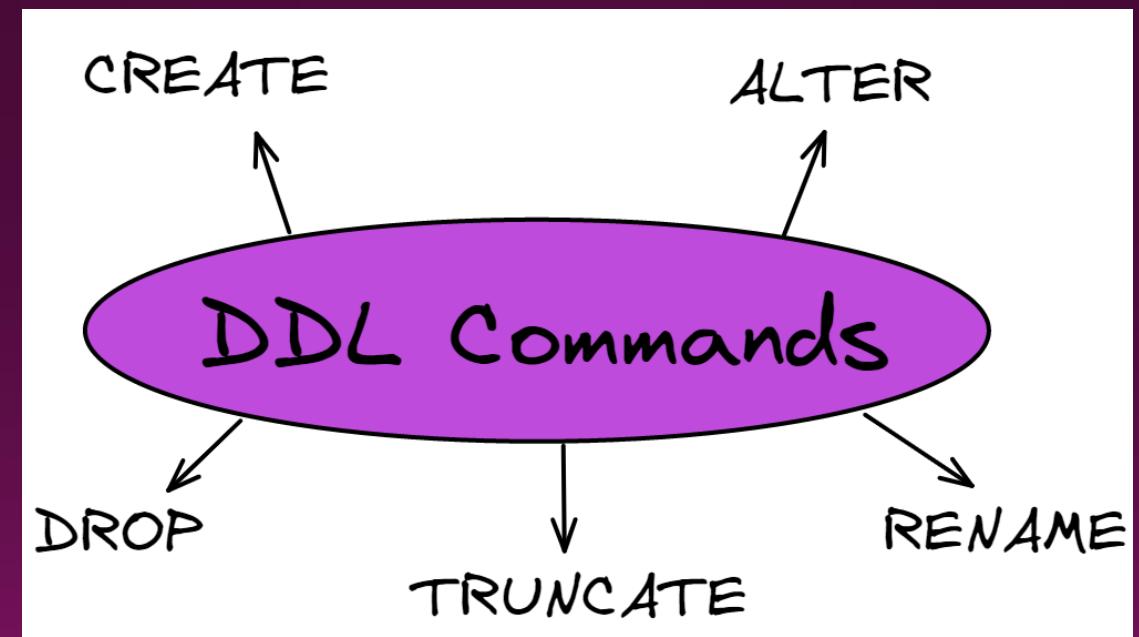


DBMS Language



Data Definition Language (DDL)

- It is used to define database structure or pattern.
- It is used to create schema, tables, indexes, constraints, etc. in the database.
- Using the DDL statements, you can create the skeleton of the database.



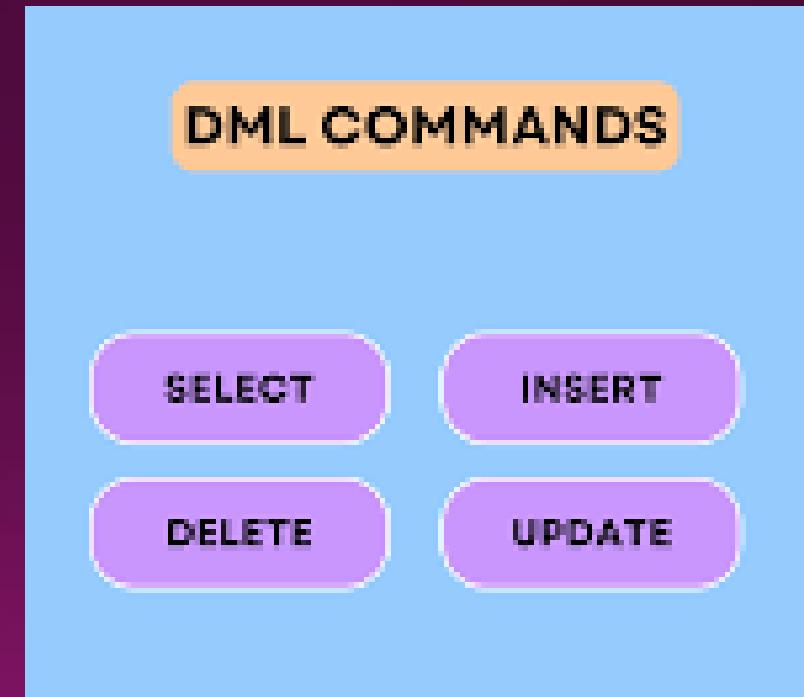
Data Definition Language (DDL)

- **Create:** It is used to create objects in the database.
- **Alter:** It is used to alter the structure of the database.
- **Drop:** It is used to delete objects from the database.
- **Truncate:** It is used to remove all records from a table.
- **Rename:** It is used to rename an object.
- **Comment:** It is used to comment on the data dictionary.

DDL
CREATE
ALTER
DROP
TRUNCATE
COMMENT
RENAME

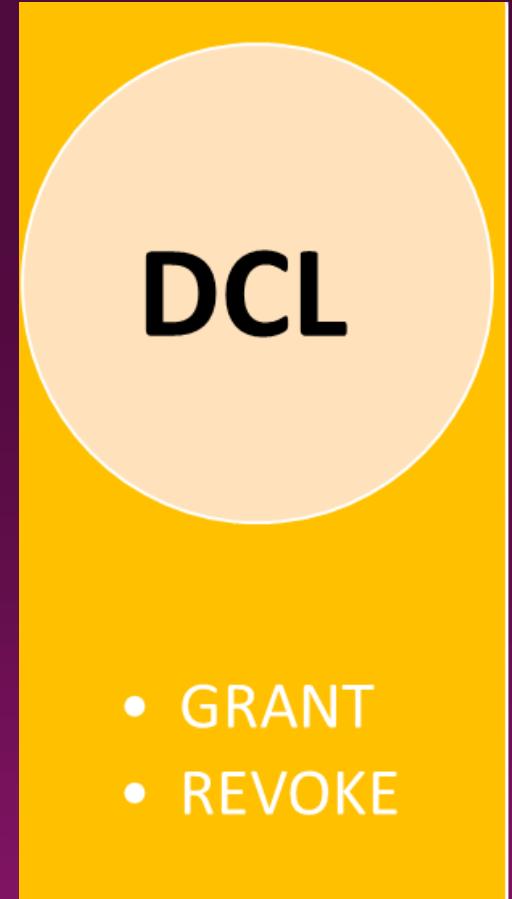
Data Manipulation Language (DML)

- It is used for accessing and manipulating data in a database.
- Here are some tasks that come under DML:
 - **Select**: It is used to retrieve data from a database.
 - **Insert**: It is used to insert data into a table.
 - **Update**: It is used to update existing data within a table.
 - **Delete**: It is used to delete all records from a table.
 - **Merge**: It performs UPSERT operation, i.e., insert or update operations.



Data Control Language (DCL)

- DCL stands for Data Control Language.
- It is used to retrieve the stored or saved data.
- The DCL execution is transactional.
- It also has rollback parameters.
- Here are some tasks that come under DCL:
 - **Grant:** It is used to give user access privileges to a database.
 - **Revoke:** It is used to take back permissions from the user.



Transaction Control Language (TCL)

- TCL is used to run the changes made by the DML statement.
- TCL can be grouped into a logical transaction.
- Here are some tasks that come under TCL:
 - Commit: It is used to save the transaction on the database.
 - Rollback: It is used to restore the database to original since the last Commit.



Key Concepts

These are fundamental components that help organize and structure data

- Tables

EmployeeID	FirstName	LastName	Position	Salary
1	John	Doe	Developer	70000
2	Jane	Smith	Manager	90000
3	Bob	Johnson	Data Analyst	60000

- Rows

- Columns

Key Concepts - Tables

- A table is a **structured representation** of data in a relational database.
- It is often compared to a spreadsheet, where data is organized into **rows and columns**.
- Tables are used to store information about a specific entity or concept, such as customers, products, or orders.

Key Concepts - Rows

- Also known as **records** or **tuples**, rows represent individual entries in a table.
- Each row contains data related to a specific instance of the entity represented by the table.

Key Concepts - Columns

- Columns, also referred to as **fields** or **attributes**, represent the different properties or attributes of the data stored in a table.
- Each column in a table has a specific data type (such as text, number, date) and holds a particular kind of information.

Terminologies used in DBMS

- Primary key
- Foreign key
- Index
- Query
- Normalization
- Transaction
- ACID Property
- Schema
- View

Introduction to Structured Query Language (SQL)

- SQL is a specialized programming language designed for managing and manipulating relational databases.
- It serves as the standard language for interacting with relational database management systems (RDBMS), allowing users to create, modify, and query databases.
- SQL provides a set of powerful and flexible commands for handling data, maintaining database integrity, and retrieving information efficiently.

Key aspects of SQL

- Data Definition Language (DDL)
- Data Manipulation Language (DML)
- Data Query Language (DQL)
- Data Control Language (DCL)
- Data Integrity
- Transactions
- Views

Data Definition Language (DDL)

- DDL defines & manage the structure of the database.
- DDL commands include:
 - **CREATE**: Used to create database objects such as tables, indexes, or views.
 - **ALTER**: Modifies the structure of existing database objects.
 - **DROP**: Deletes database objects like tables or indexes.

Data Manipulation Language (DML)

- DML statements are used to manipulate the data stored in the database.
 - **SELECT**: Retrieves data from one or more tables.
 - **INSERT**: Adds new records to a table.
 - **UPDATE**: Modifies existing records in a table.
 - **DELETE**: Removes records from a table.

Data Query Language (DQL)

- DQL is a subset of SQL used exclusively for querying and retrieving data from the database.
- The primary DQL command is **SELECT**.

Data Control Language (DCL)

DCL statements are used to control access to data within the database.

- **GRANT**: Provides specific privileges to database users.
- **REVOKE**: Removes specific privileges from database users.

Data Integrity

- SQL allows the definition of constraints to maintain data integrity.
 - **PRIMARY KEY**: Ensures the uniqueness of a column's values in a table.
 - **FOREIGN KEY**: Establishes a link between two tables based on a column.
 - **CHECK**: Defines conditions that must be satisfied for data to be entered into a table.
 - **UNIQUE**: Ensures that all values in a column are unique.

Transactions

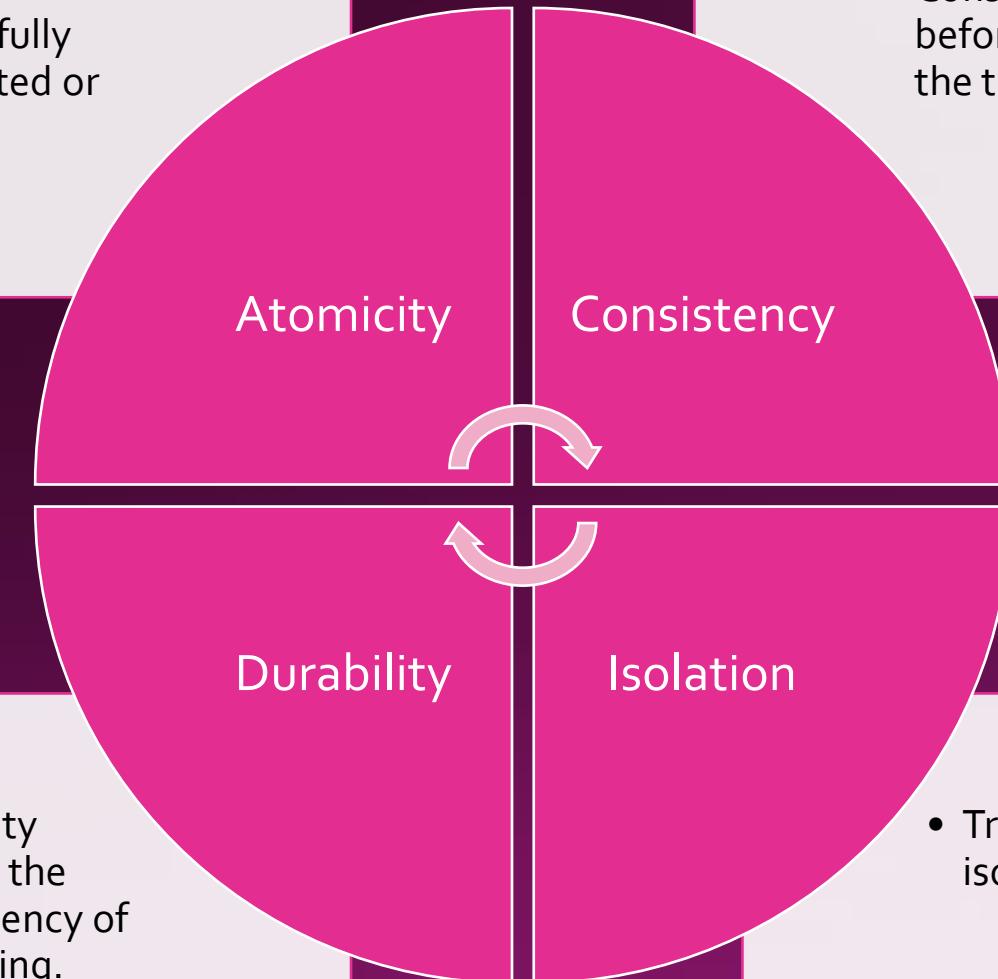
- SQL supports transaction management to ensure the consistency and reliability of database operations.
- The **ACID** properties (Atomicity, Consistency, Isolation, Durability) define the characteristics of a transaction.

Views

- SQL allows the creation of virtual tables known as views.
- Views are based on the result of a SELECT query and provide a way to simplify complex queries or restrict access to certain columns.

ACID Properties

- Either successfully completed or none.



- Consistent state before & after the transaction.

- Durability ensures the permanency of something.

- Transactions are isolation

SQL

- It's a language.
- Used for interacting with Relational DBMS (RDBMS).
- Performs C.R.U.D operation.
 - C = Create
 - R = Read/Retrieve
 - U = Update
 - D = Delete
- Also performs admin tasks:
 - Security
 - User management
 - Import/Export data
 - Backup

Create Database

Syntax to create a DB:

- CREATE DATABASE <DB-NAME>;

```
create database demo01;
```

Selecting DB

- In simple terms, the use statement selects a specific database and then performs operations on it using the inbuilt commands of SQL.

Syntax:

- USE DATABASE;

```
use demo01;
```

Creating Table

```
CREATE TABLE employees (
    employee_id INT PRIMARY KEY,
    first_name VARCHAR(50),
    last_name VARCHAR(50),
    salary DECIMAL(10, 2)
);
```

Viewing table

```
select * from employees;
```

Output:

employee_id	first_name	last_name	salary
NULL	NULL	NULL	NULL

Inserting 1st row in table

```
INSERT INTO employees (
employee_id,
first_name,
last_name,
salary)
VALUES (1, 'John', 'Doe', 50000);
```

Verifying:

```
select * from employees;
```

Output:

	employee_id	first_name	last_name	salary
	1	John	Doe	50000.00
	NULL	NULL	NULL	NULL

Inserting multiple values

```
# Inserting multiple rows
INSERT INTO employees (employee_id, first_name,
last_name, salary) VALUES
(2, 'User1', 'one', 30000),
(3, 'User2', 'two', 40000),
(4, 'User3', 'three', 25000),
(5, 'User4', 'four', 32000);
```

Verifying:

```
select * from employees;
```

Output:

employee_id	first_name	last_name	salary
1	John	Doe	50000.00
2	User1	one	30000.00
3	User2	two	40000.00
4	User3	three	25000.00
5	User4	four	32000.00

Dropping

- Table
- Database

```
drop table employees;
```

```
drop database demo01;
```

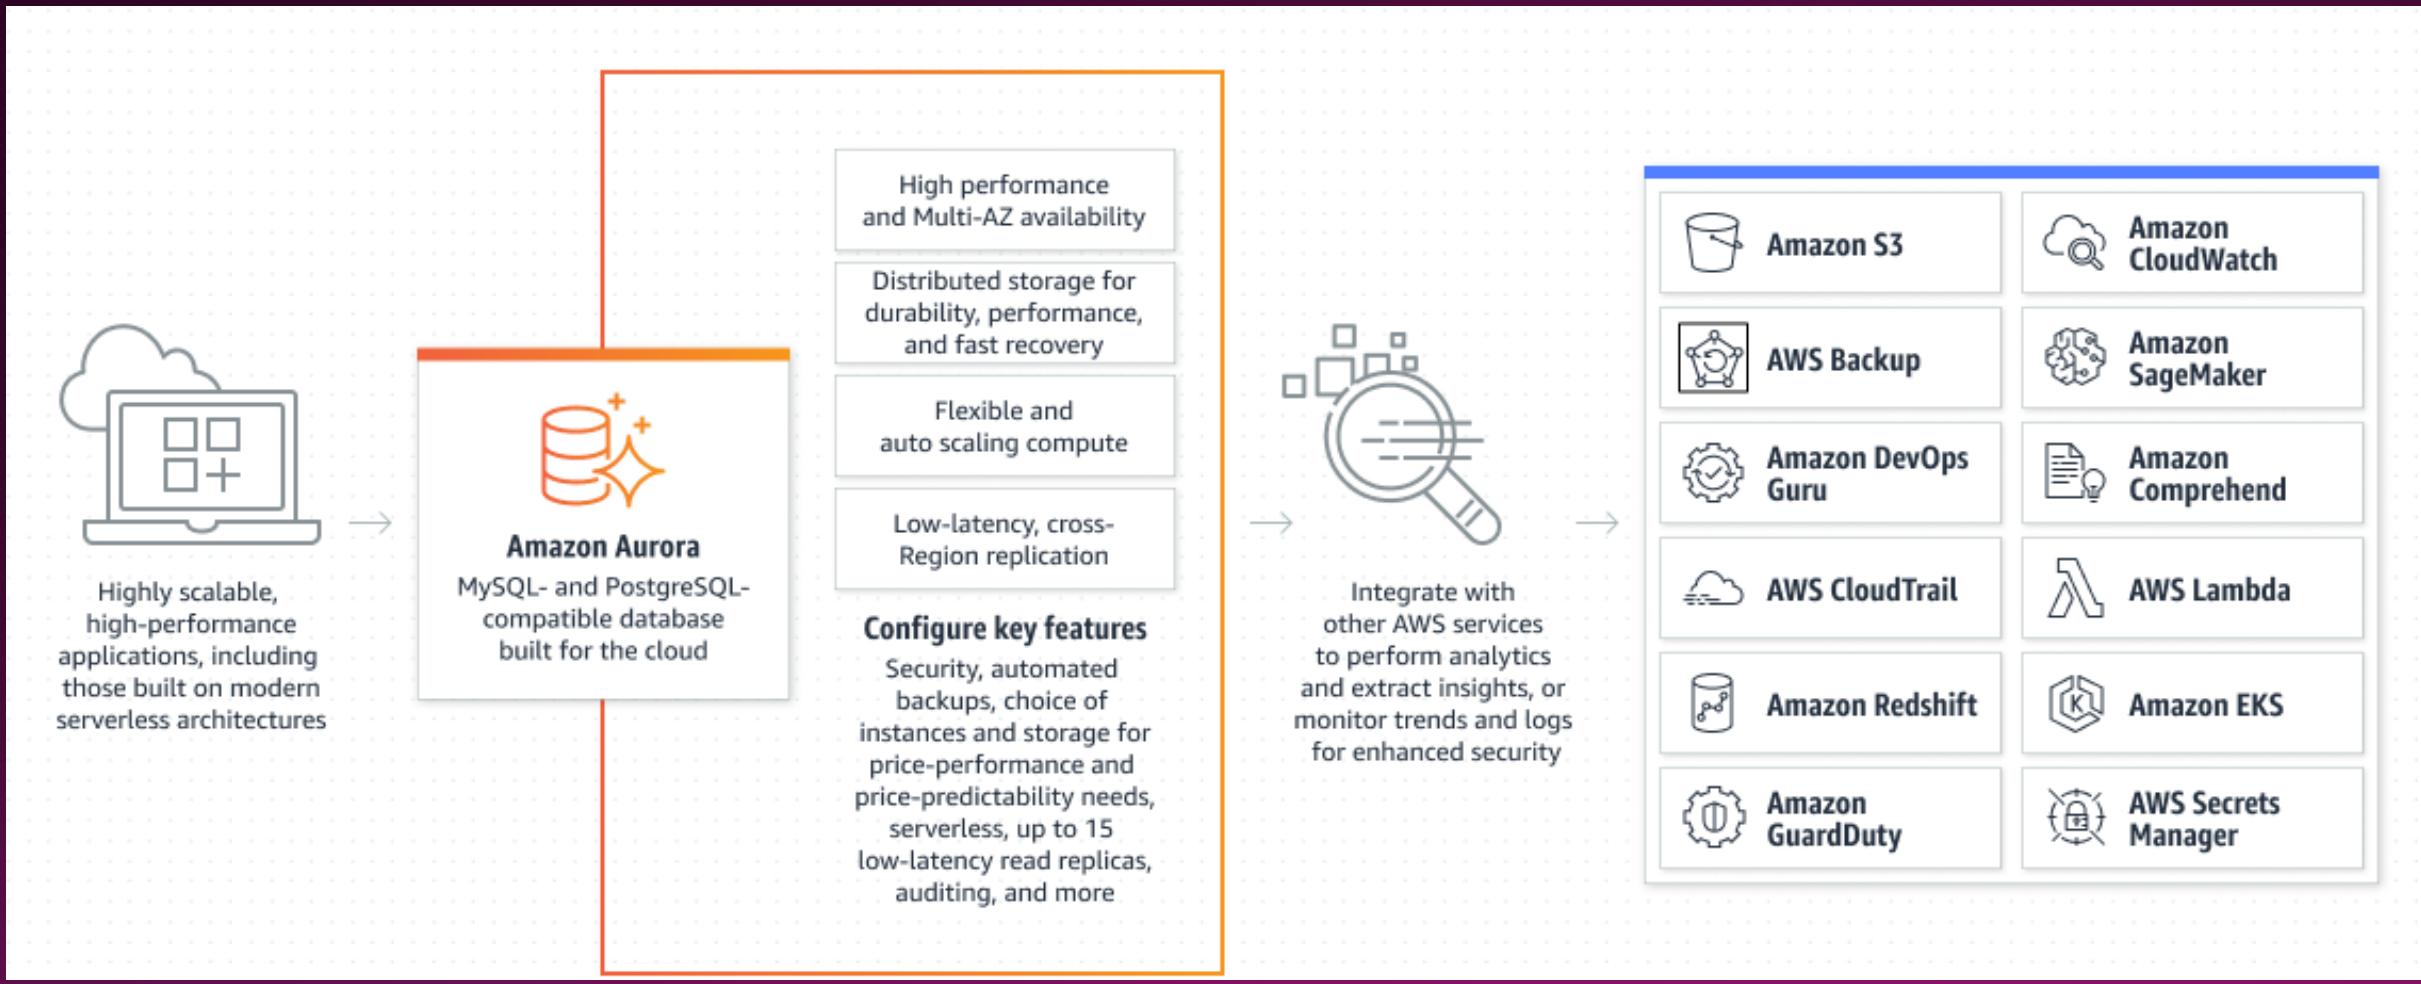
12 | 16:30:39 | drop table employees

14 | 16:32:57 | drop database demo01

AWS Aurora

- Amazon Aurora is a fully managed relational database service provided by Amazon Web Services (AWS).
- It is compatible with MySQL and PostgreSQL, offering high performance, availability, and durability.
- Amazon Aurora is designed to provide the benefits of commercial databases with the cost-effectiveness and simplicity of open-source databases.

AWS Aurora - Architecture



AWS Aurora – What it does?

- Amazon Aurora provides a fully managed, highly available, and scalable relational database engine.
- It is designed to deliver high performance and reliability while reducing the operational overhead typically associated with managing traditional relational databases.
- Aurora uses a distributed and fault-tolerant architecture to ensure high availability and durability of data.

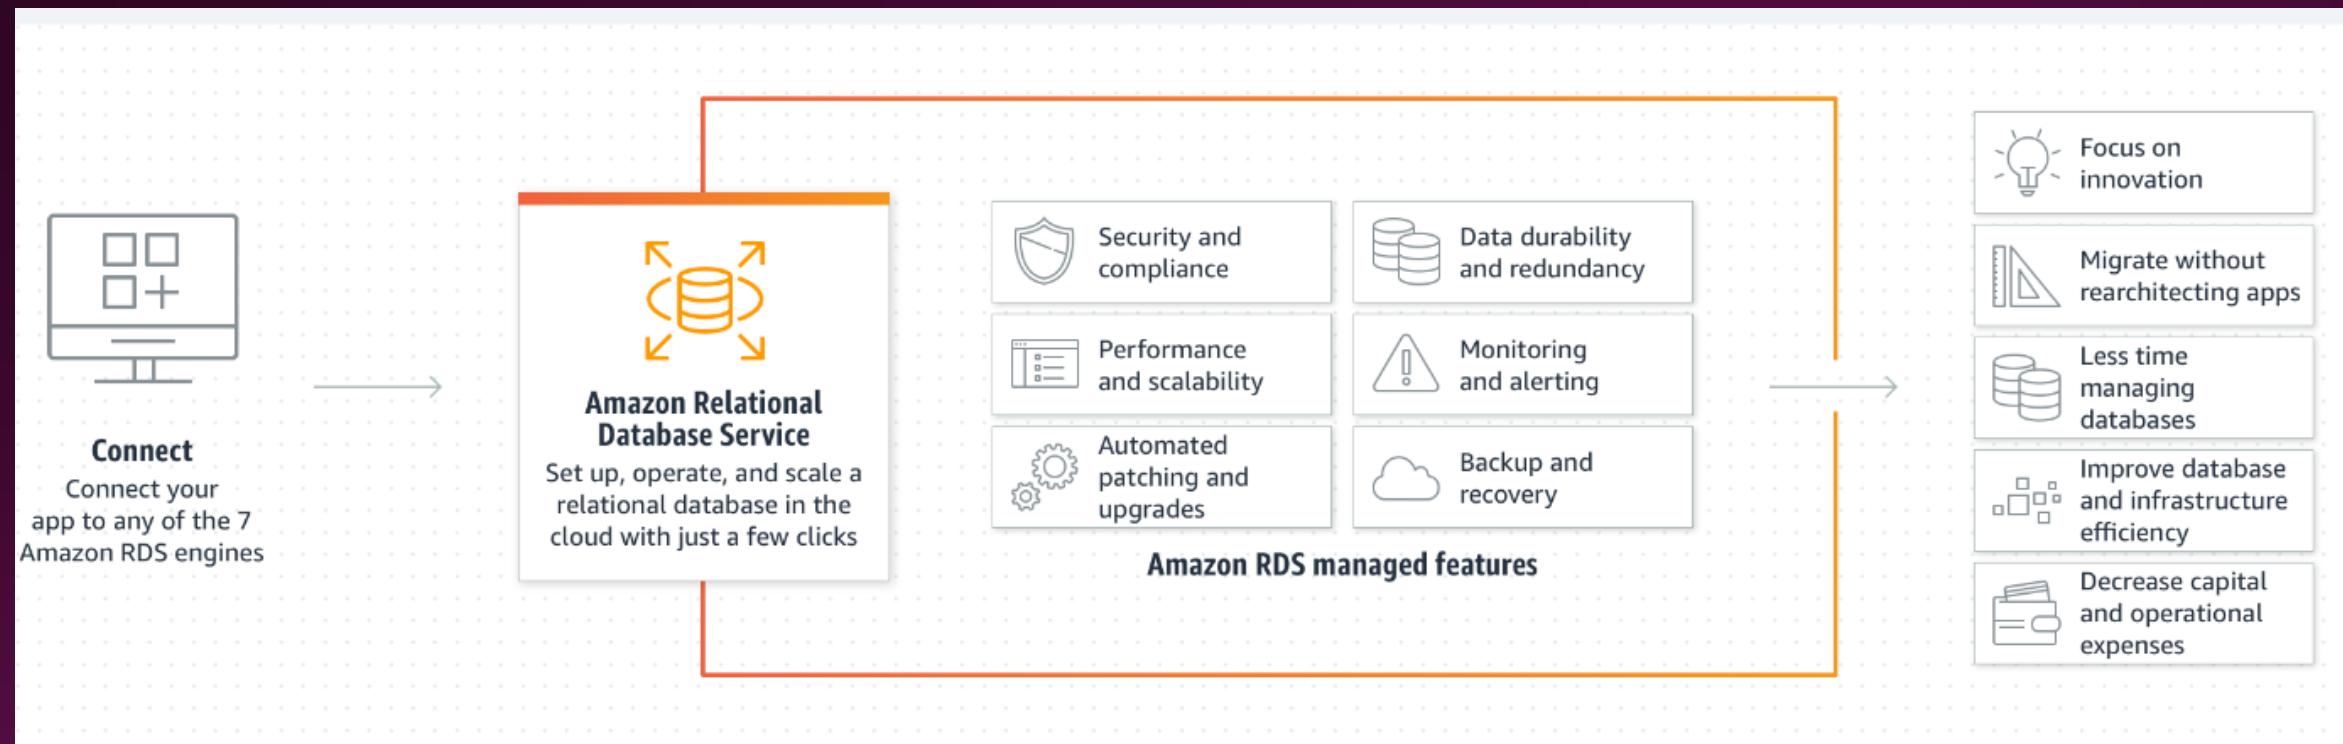
AWS Aurora – Key Features

- High Availability
- Performance
- Auto-Scaling
- Global Databases
- Backup & Snapshots
- Security
- Read Replicas

Amazon Relation Database Service (RDS)

- Amazon RDS is a fully managed relational database service provided by Amazon Web Services (AWS).
- It enables users to set up, operate, and scale relational databases in the cloud without the need for manual intervention in common database administration tasks.
- RDS supports several popular relational database engines, making it a versatile and scalable solution for a wide range of applications.

Amazon RDS - Architecture



Amazon RDS - What it does?

- Amazon RDS automates time-consuming administrative tasks such as hardware provisioning, database setup, patching, and backups.
- It allows users to focus on building and optimizing their applications rather than managing the underlying infrastructure.
- RDS supports multiple database engines, including MySQL, PostgreSQL, MariaDB, Oracle, and Microsoft SQL Server.

Amazon RDS – Use Cases

- Web and Mobile Applications
- E-commerce Applications
- Content Management Systems (CMS)
- Data Warehousing
- Business Applications
- Dev/Test Environments
- WordPress and Content-Based Websites

Database Engines

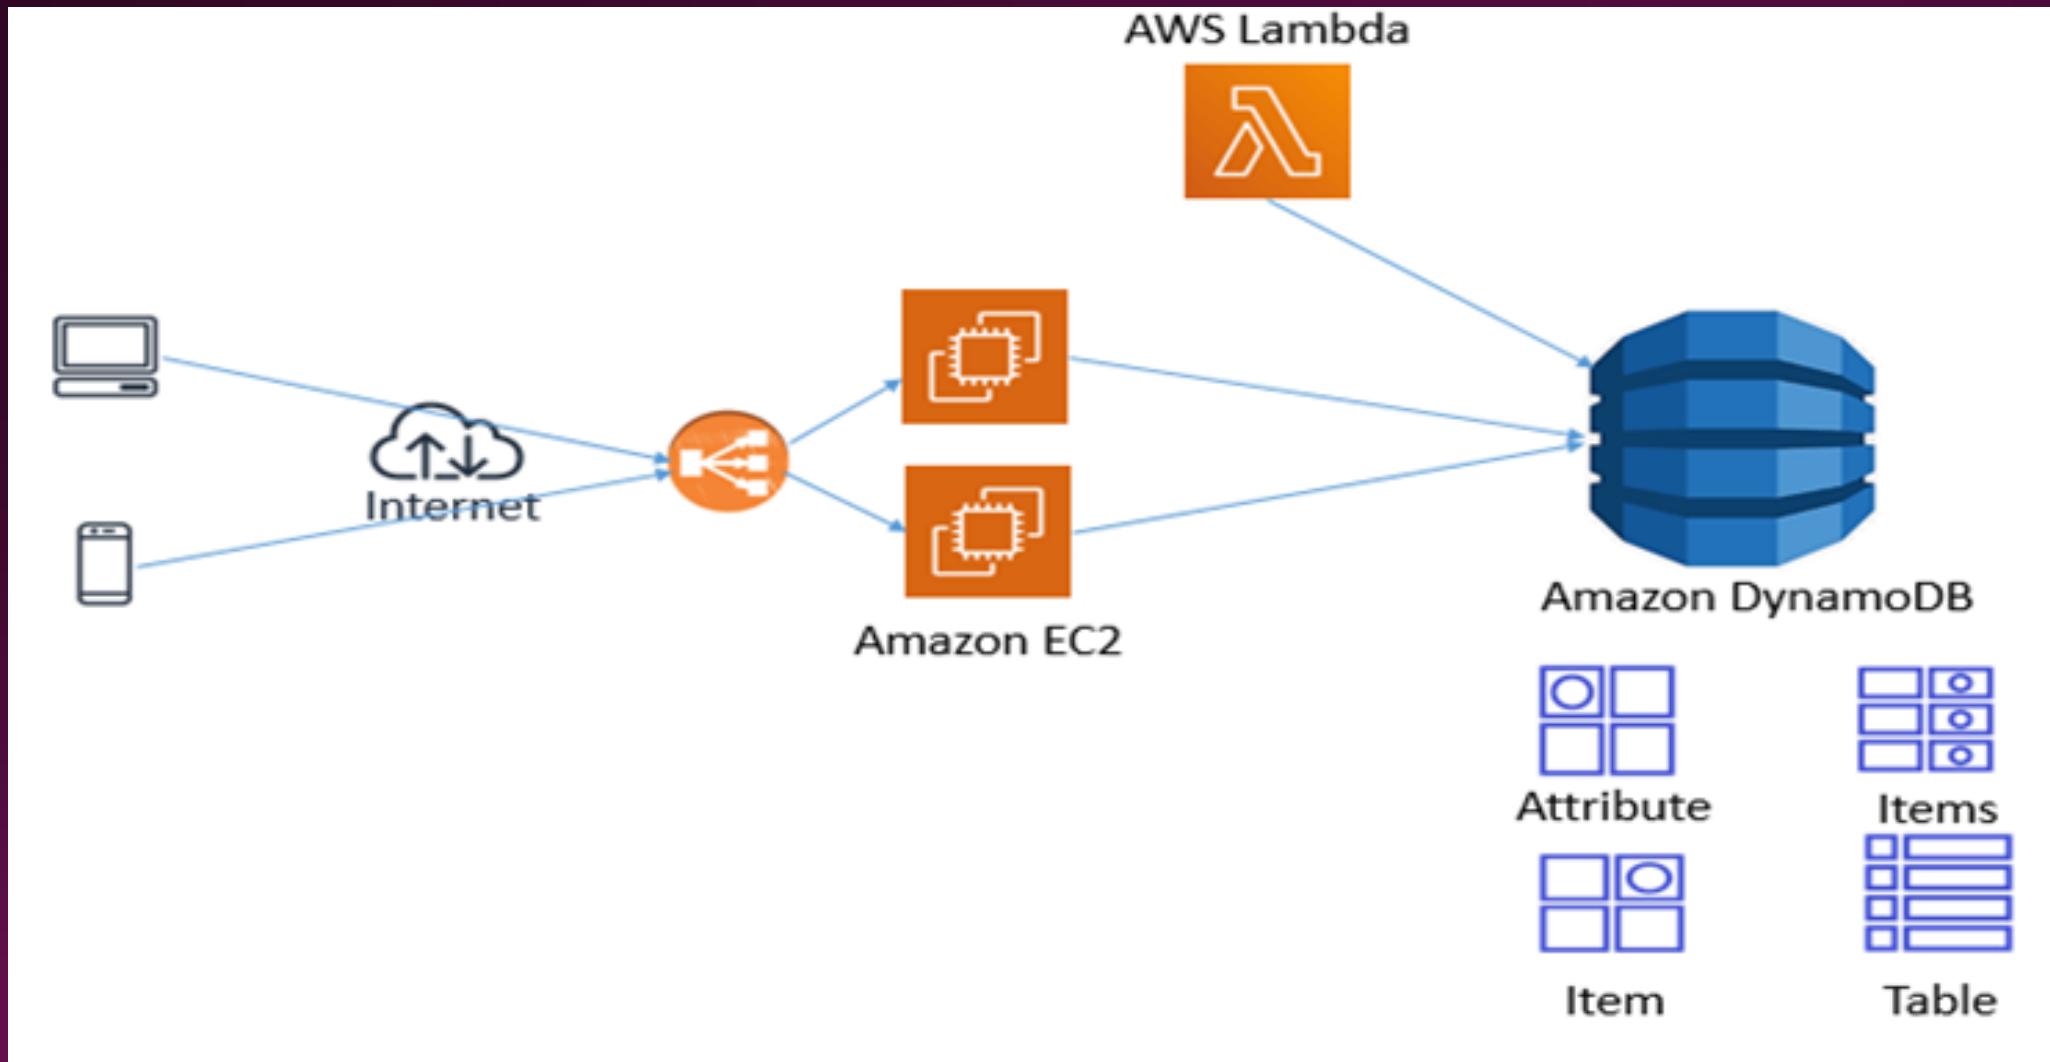
- MySQL
- PostgreSQL
- MariaDB
- Oracle
- Microsoft SQL Server



Amazon DynamoDB

- Amazon DynamoDB is a fully managed NoSQL database service provided by Amazon Web Services (AWS).
- It is designed to deliver high performance, scalability, and low-latency access to applications that require fast and predictable performance at any scale.
- DynamoDB is a serverless database, which means AWS takes care of the operational aspects, allowing developers to focus on building applications.

Amazon DynamoDB



Amazon DynamoDB – What it does?

- DynamoDB is a key-value and document database that provides single-digit millisecond latency at any scale.
- It supports both document and key-value data models, and its architecture is optimized for horizontal scalability.
- DynamoDB automatically replicates data across multiple Availability Zones within an AWS region to ensure high availability and fault tolerance.

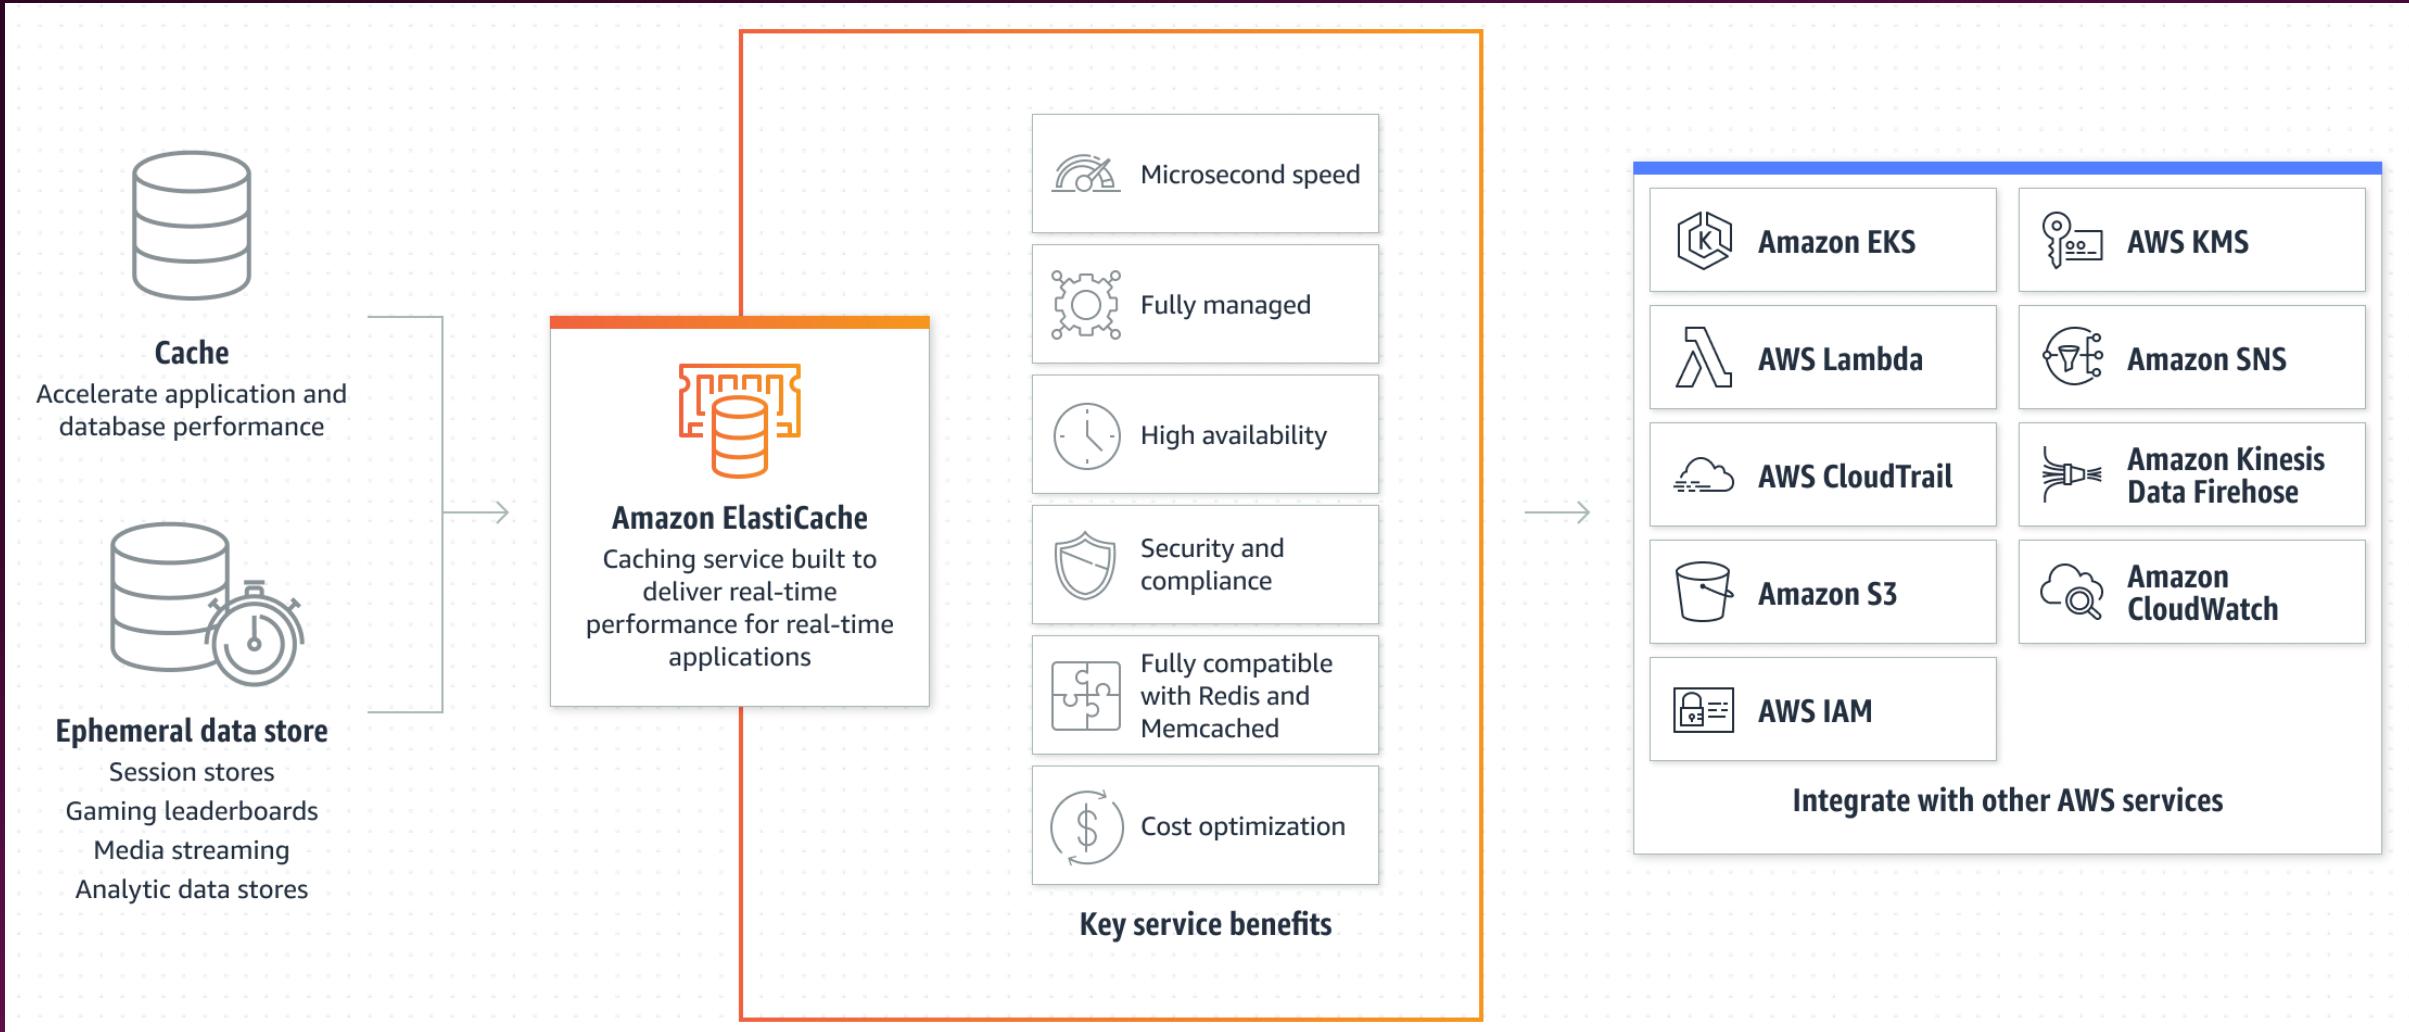
Amazon DynamoDB – Use Cases

- Web and Mobile Applications
- Gaming Applications
- IoT (Internet of Things)
- Ad Tech
- Caching and Metadata Storage
- E-commerce Platforms
- Content Management Systems (CMS)
- Real-Time Analytics
- Session Management
- Caching and Metadata Storage

AWS ElastiCache

- Amazon ElastiCache is a fully managed, in-memory caching service provided by AWS.
- It is designed to improve the performance of web applications by allowing them to retrieve data from a fast, in-memory cache rather than relying on slower, disk-based databases.
- ElastiCache supports two popular open-source in-memory caching engines:
 - Redis and
 - Memcached.

AWS ElastiCache - Architecture



AWS ElastiCache - What it does?

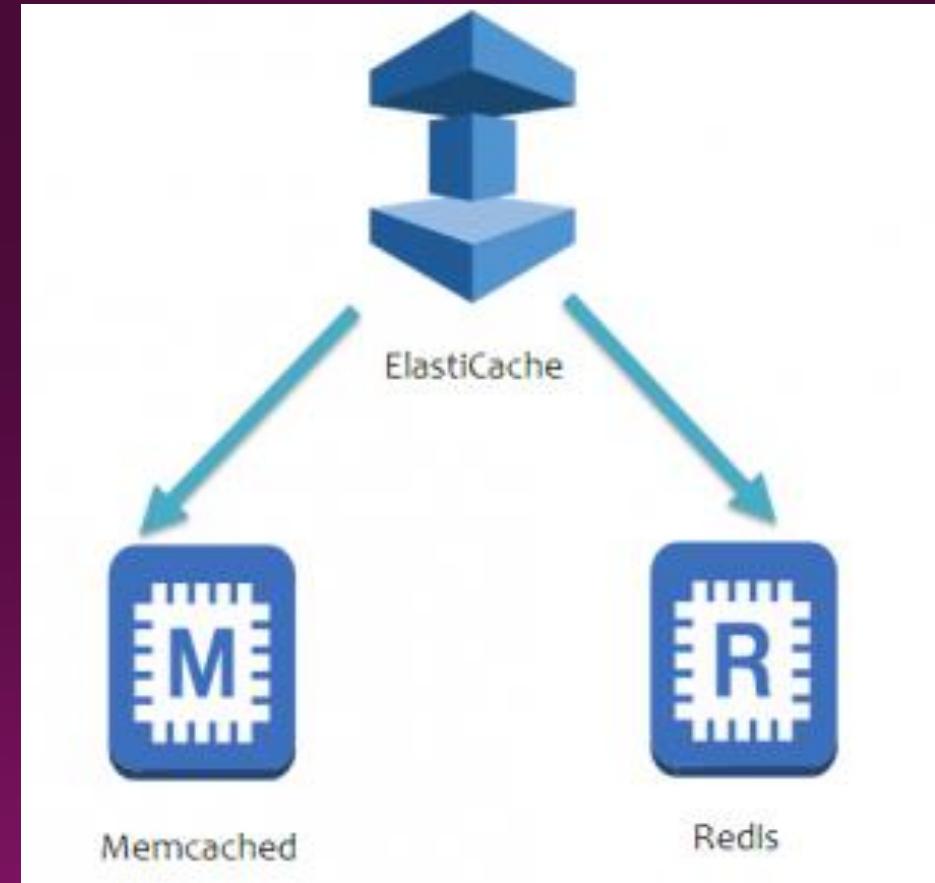
- ElastiCache enhances the performance and scalability of applications by providing a fully managed, in-memory caching layer.
- It stores frequently accessed data in-memory, reducing the need for repeated queries to the underlying databases.
- This results in faster response times and improved overall application performance.

AWS ElastiCache – Use Cases

- Caching Frequently Accessed Data
- Session Store
- Real-Time Analytics
- Leaderboards and Counting
- Pub/Sub Messaging
- Geospatial Data
- Content Caching
- Application State Management

AWS ElastiCache – Key Features

- Managed Service
- Scalability
- High Availability
- Security
- Monitoring and Logging
- Automatic Backups
- Parameter Groups



Database Design and Modelling

- Modeling simplifies database design and maintenance by enabling you, the data architect, to visualize requirements and resolve design issues.
- Model-driven database design is an efficient methodology for creating valid and well-performing databases, while providing the flexibility to respond to evolving data requirements.
- Models are used to build EER diagrams and physical MySQL databases.

What is good database design?

- Divides your information into subject-based tables to reduce redundant data.
- Provides Access with the information it requires to join the information in the tables together as needed.
- Helps support and ensure the accuracy and integrity of your information.
- Accommodates your data processing and reporting needs.

The design process

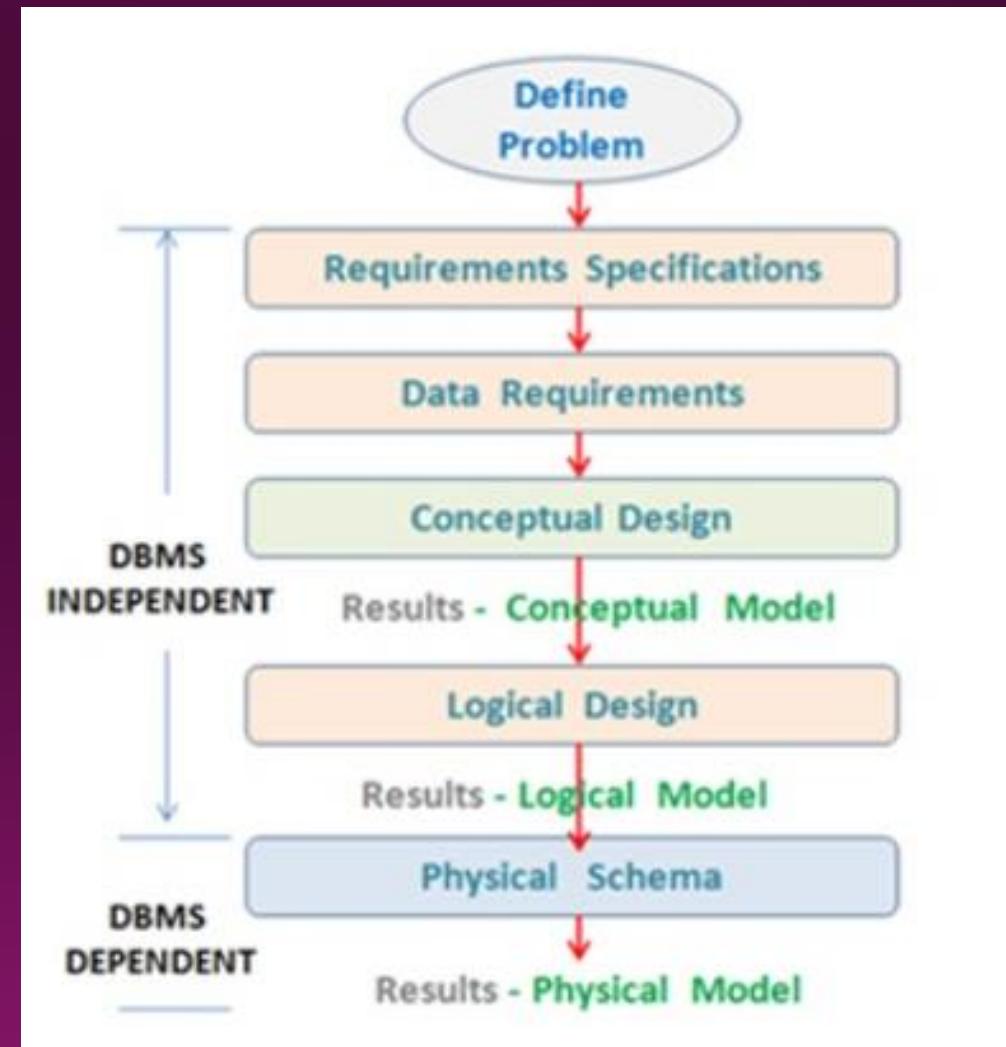
Step1: Defining the Database Objectives

Step2: Database design team & stake holders

Step3: Mapping business processes, rule & policies.

Step4: Defining the user data requirements

Step5: Creating data models



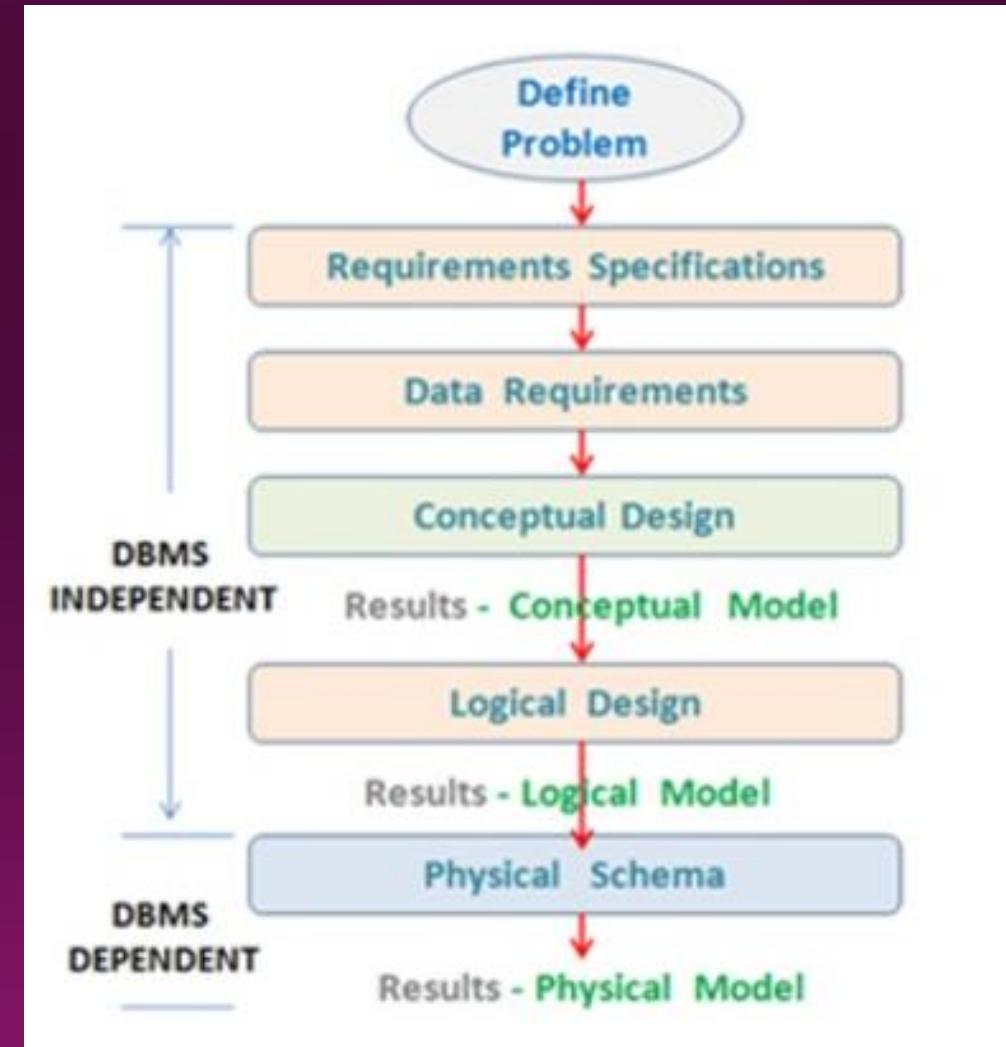
The design process

Step6: Creating conceptual data model

Step7: Creating Logical data model

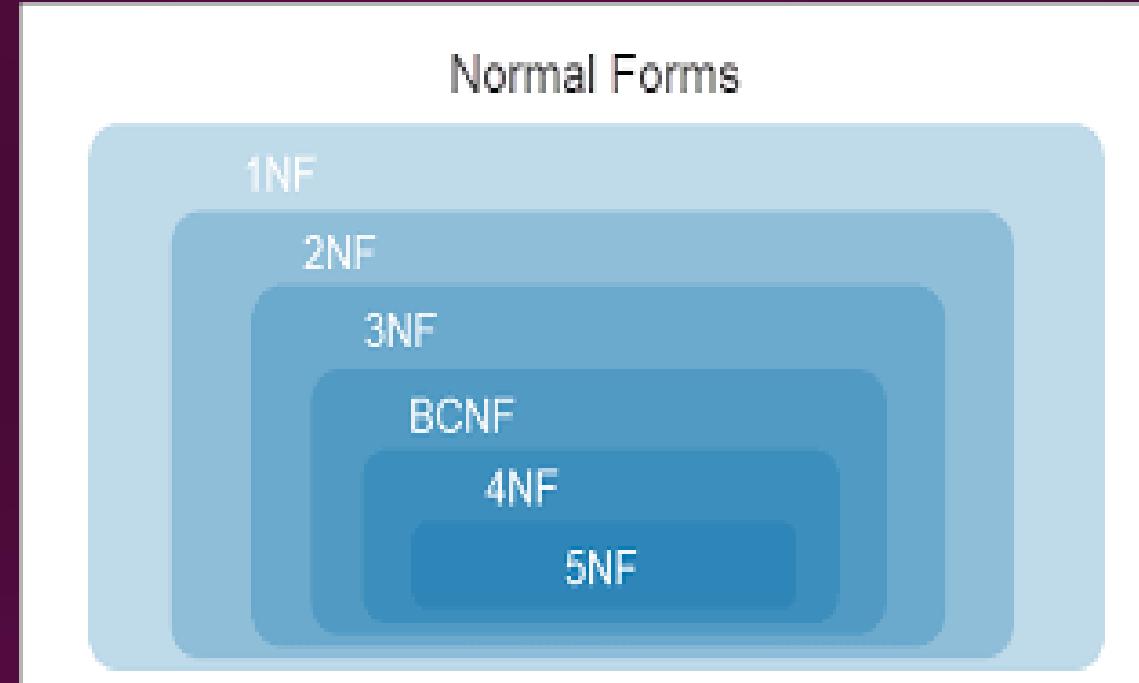
Step8: Database normalization

Step9: Defining the database requirements
specifications



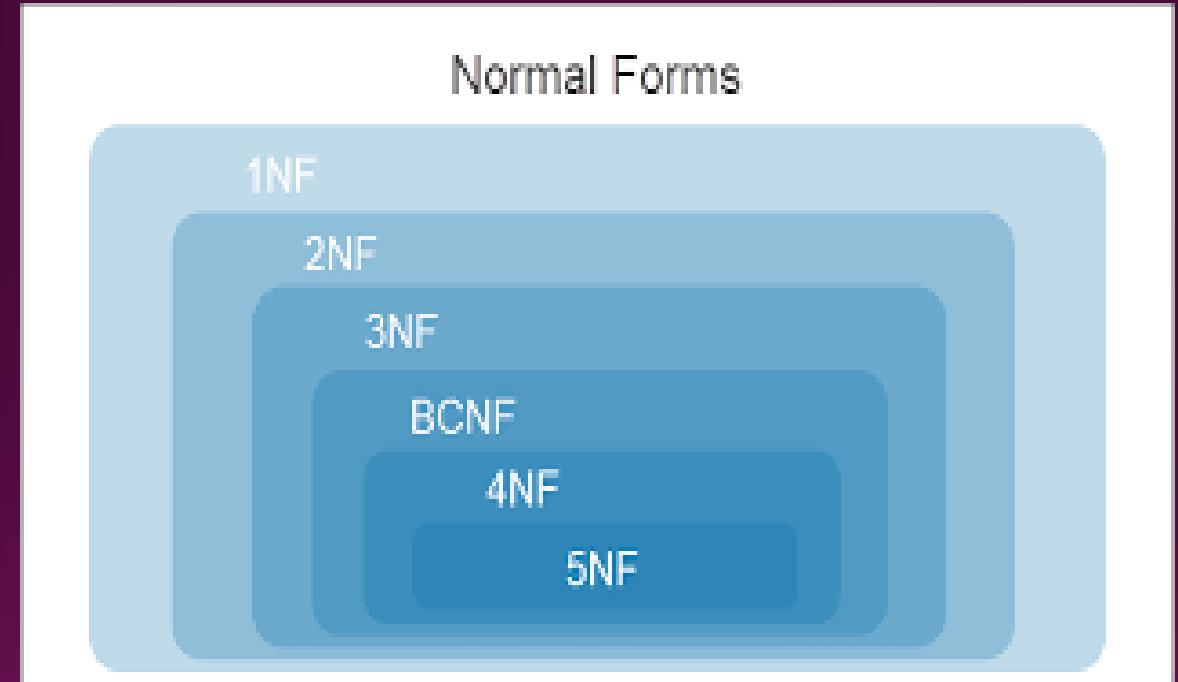
Normalization

- Normalization is the process of minimizing redundancy from a relation or set of relations.
- Redundancy in relation may cause insertion, deletion, and update anomalies. So, it helps to minimize the redundancy in relations.
- Normal forms are used to eliminate or reduce redundancy in database tables.



Normalization

- First Normal Form (1NF)
- Second Normal Form (2NF)
- Third Normal Form (3NF)
- Boyce-Codd Normal Form (BCNF)
- Fourth Normal Form (4NF)
- Fifth Normal Form (5NF)



Normalization of DBMS – 1NF

- This is the most basic level of normalization.
- In 1NF, each table cell should contain only a single value, and each column should have a unique name.
- The first normal form helps to eliminate duplicate data and simplify queries.

Normalization of DBMS – 2NF

- 2NF eliminates redundant data by requiring that each non-key attribute be dependent on the primary key.
- This means that each column should be directly related to the primary key, and not to other columns.

Normalization of DBMS – 3NF

- 3NF builds on 2NF by requiring that all non-key attributes are independent of each other.
- This means that each column should be directly related to the primary key, and not to any other columns in the same table.

Boyce-Codd Normal Form (BCNF):

- BCNF is a stricter form of 3NF that ensures that each determinant in a table is a candidate key.
- In other words, BCNF ensures that each non-key attribute is dependent only on the candidate key.